

Structural Causal Bandits with non-manipulable variables

Sanghack Lee Elias Bareinboim

*We only recently realized about preparing draft presentation
so there might be some missing pieces (e.g., animation).



Executive Summary

- **SCM-MAB** = MAB + Causality

Executive Summary

- **SCM-MAB** = MAB + Causality
where actions = interventions

Executive Summary

- **SCM-MAB** = MAB + Causality
where actions = interventions
- Q: Which interventions *can* be optimal?

Executive Summary

- **SCM-MAB** = MAB + Causality
where actions = interventions
- Q: Which interventions *can* be optimal?
A: Possibly-Optimal Minimal Intervention Set (**POMIS**)

Executive Summary

- **SCM-MAB** = MAB + Causality
where actions = interventions
- Q: Which interventions *can* be optimal?
A: Possibly-Optimal Minimal Intervention Set (**POMIS**)
- Q: How to learn *an arm's* reward from *other arms*?

Executive Summary

- **SCM-MAB** = MAB + Causality
where actions = interventions
- Q: Which interventions *can* be optimal?
A: Possibly-Optimal Minimal Intervention Set (**POMIS**)
- Q: How to learn *an arm's* reward from *other arms*?
A: a generalized identifiability algorithm (**z²ID**)

Executive Summary

- **SCM-MAB** = MAB + Causality
where actions = interventions
- Q: Which interventions *can* be optimal?
A: Possibly-Optimal Minimal Intervention Set (**POMIS**)
- Q: How to learn *an arm's* reward from *other arms*?
A: a generalized identifiability algorithm (**z²ID**)
- Q: How to utilize **POMIS** and **z²ID** in bandit algorithm?

Executive Summary

- **SCM-MAB** = MAB + Causality
where actions = interventions
- Q: Which interventions *can* be optimal?
A: Possibly-Optimal Minimal Intervention Set (**POMIS**)
- Q: How to learn *an arm's* reward from *other arms*?
A: a generalized identifiability algorithm (**z²ID**)
- Q: How to utilize **POMIS** and **z²ID** in bandit algorithm?
A: modified MAB algorithms for SCM-MAB (**z²-TS**, **z²-kl-UCB**)

Executive Summary

- **SCM-MAB** = MAB + Causality
where actions = interventions
- Q: Which interventions *can* be optimal?
A: Possibly-Optimal Minimal Intervention Set (**POMIS**)
- Q: How to learn *an arm's* reward from *other arms*?
A: a generalized identifiability algorithm (**z²ID**)
- Q: How to utilize **POMIS** and **z²ID** in bandit algorithm?
A: modified MAB algorithms for SCM-MAB (**z²-TS**, **z²-kl-UCB**)
- Faster convergence: smaller # of arms; more accurate estimation.

Overview

- **Motivation:** why we need to be causally-sensible

Overview

- **Motivation:** why we need to be causally-sensible
- **SCM-MAB** and its **structural properties**

Overview

- **Motivation:** why we need to be causally-sensible
- **SCM-MAB** and its **structural properties**
- **SCM-MAB algorithms**

Overview

- **Motivation:** why we need to be causally-sensible
- **SCM-MAB** and its **structural properties**
- **SCM-MAB algorithms**
- **Empirical results**

Overview

- **Motivation:** why we need to be causally-sensible
- **SCM-MAB** and its **structural properties**
- **SCM-MAB algorithms**
- **Empirical results**
- **Conclusions**

Motivation

Multi-armed bandit (MAB)

A classic, sequential decision-making problem

- Given: a set of **arms** (actions), \mathbf{A}
- How: at round t , pull an arm A_t , and get a **reward** Y_{A_t}
- Goal: to minimize cumulative **regret** (or maximize cumulative reward)


a trade-off between **exploitation** vs. **exploration**

Examples: ad. placement, online news recommendation, packet routing ...

Key **assumption**: arms are *independent* (in a traditional MAB setting)

Multi-armed bandit (MAB)

The reward mechanism can be understood as (at its simplest form possible),



Can we be agnostic to the mechanism between A and Y ?

What if there exists a complex (causal) mechanism?

MAB through Causal Lens

Structural Causal Model (SCM) $\mathcal{M} = \langle \mathbf{U}, \mathbf{V}, \mathbf{F}, P(\mathbf{U}) \rangle$:


- \mathbf{U} : unobserved variables
- \mathbf{V} : observed variables
- \mathbf{F} : a set of functions for \mathbf{V}
- $P(\mathbf{U})$: a joint distribution over \mathbf{U} (\sim randomness)

MAB through Causal Lens

Structural Causal Model (SCM) $\mathcal{M} = \langle \mathbf{U}, \mathbf{V}, \mathbf{F}, P(\mathbf{U}) \rangle$:

- \mathbf{U} : unobserved variables
- \mathbf{V} : observed variables
- \mathbf{F} : a set of functions for \mathbf{V}
- $P(\mathbf{U})$: a joint distribution over \mathbf{U} (\sim randomness)

A causal graph \mathcal{G} conforming to \mathcal{M} looks like **DAG** + **bidirected edges** for unobserved confounders (UCs).¹




¹among \mathbf{U} , only UCs will be visualized.

MAB through Causal Lens

Structural Causal Model (SCM) $\mathcal{M} = \langle \mathbf{U}, \mathbf{V}, \mathbf{F}, P(\mathbf{U}) \rangle$:

- \mathbf{U} : unobserved variables
- \mathbf{V} : observed variables
- \mathbf{F} : a set of functions for \mathbf{V}
- $P(\mathbf{U})$: a joint distribution over \mathbf{U} (\sim randomness)


A causal graph \mathcal{G} conforming to \mathcal{M} looks like **DAG** + **bidirected edges** for unobserved confounders (UCs).¹



¹among \mathbf{U} , only UCs will be visualized.

MAB through Causal Lens

An example with a traditional MAB problem



- a bandit algorithm plays an arm a by *doing* $do(a)$,
- get a reward, e.g.,


$$Y = f(a, u) = \mu_a + u,$$

where, e.g., $U \sim \mathcal{N}(0, 1)$.

(with time step t implicit)


MAB through Causal Lens — a worst case scenario?

Given an underlying causal mechanism,




MAB through Causal Lens — a worst case scenario?

Given an underlying causal mechanism,




Ignorant to such causal mechanism,




MAB through Causal Lens — a worst case scenario?

Given an underlying causal mechanism,





Ignorant to such causal mechanism,



Insensitive to the structure: $\mathbf{A} = \mathfrak{X}_X \times \mathfrak{X}_Z$ (simultaneously)

Sensitive to the structure: $\mathbf{A} = \bigcup_{\mathbf{w} \subseteq \{X,Z\}} \mathfrak{X}_{\mathbf{w}}$ (all combinations)


MAB through Causal Lens — a worst case scenario?



Insensitive to the structure: $\mathbf{A} = \mathfrak{X}_X \times \mathfrak{X}_Z$ (simultaneously)

Sensitive to the structure: $\mathbf{A} = \bigcup_{\mathbf{w} \subseteq \{X, Z\}} \mathfrak{X}_{\mathbf{w}}$ (all combinations)

MAB through Causal Lens — a worst case scenario?



$\mathcal{M} = \langle \{U_X, U_Y, U_Z, U_{YZ}\}, \{X, Y, Z\}, \mathbf{F}, P(\mathbf{U}) \rangle$ where \mathbf{F} is

$$X \leftarrow U_X$$

$$Z \leftarrow U_Z \oplus X \oplus U_{YZ}$$

$$Y \leftarrow U_Y \oplus Z \oplus U_{YZ}$$

and $P(U_X = 1) = 0.6$, $P(U_Y = 1) = 0.15$, $P(U_Z = 1) = 0.11$,
 $P(U_{YZ} = 1) = 0.51$.

MAB through Causal Lens — a worst case scenario?

Can we do better than 'all subsets' approach if we are aware of the underlying causal graph?

SCM-MAB

SCM-MAB, definition

A **SCM-MAB** is $\langle M, Y, \mathbf{N} \rangle$:

- a SCM $\mathcal{M} = \langle \mathbf{U}, \mathbf{V}, \mathbf{F}, P(\mathbf{U}) \rangle$,
- a reward variable $Y \in \mathbf{V}$,
- non-manipulable variables $\mathbf{N} \subseteq \mathbf{V} \setminus \{Y\}$

SCM-MAB, definition

A **SCM-MAB** is $\langle M, Y, \mathbf{N} \rangle$:

- a SCM $\mathcal{M} = \langle \mathbf{U}, \mathbf{V}, \mathbf{F}, P(\mathbf{U}) \rangle$,
- a reward variable $Y \in \mathbf{V}$,
- non-manipulable variables $\mathbf{N} \subseteq \mathbf{V} \setminus \{Y\}$

Therefore,

- Actions: $\mathbf{A} = \{\mathbf{x} \in \mathfrak{X}_{\mathbf{X}} | \mathbf{X} \subseteq \mathbf{V} \setminus \mathbf{N} \setminus \{Y\}\}$ (including observation)
- Reward distribution: $P(Y | do(\mathbf{X} = \mathbf{x}))$ (or $P_{\mathbf{x}}(Y)$) ($\forall \mathbf{x} \in \mathbf{A}$)
- Expected reward: $\mu_{\mathbf{x}} = \mathbb{E}[Y | do(\mathbf{x})]$

SCM-MAB, definition

A **SCM-MAB** is $\langle M, Y, \mathbf{N} \rangle$:


- a SCM $\mathcal{M} = \langle \mathbf{U}, \mathbf{V}, \mathbf{F}, P(\mathbf{U}) \rangle$,
- a reward variable $Y \in \mathbf{V}$,
- non-manipulable variables $\mathbf{N} \subseteq \mathbf{V} \setminus \{Y\}$

Therefore,

- Actions: $\mathbf{A} = \{\mathbf{x} \in \mathfrak{X}_{\mathbf{X}} | \mathbf{X} \subseteq \mathbf{V} \setminus \mathbf{N} \setminus \{Y\}\}$ (including observation)
- Reward distribution: $P(Y | do(\mathbf{X} = \mathbf{x}))$ (or $P_{\mathbf{x}}(Y)$) ($\forall \mathbf{x} \in \mathbf{A}$)
- Expected reward: $\mu_{\mathbf{x}} = \mathbb{E}[Y | do(\mathbf{x})]$


Assumption: we access to the causal graph \mathcal{G} without knowing \mathbf{F} nor $P(\mathbf{U})$.

SCM-MAB examples




- No non-manipulable variable
- Intervention sets: \emptyset , $\{X\}$, $\{Z\}$, $\{X, Z\}$
- Arms: $do(\emptyset)$, $do(X = 0)$, $do(X = 1)$, \dots , $do(X = 1, Z = 1)$

SCM-MAB examples



- Z is non-manipulable
- Intervention sets: $\emptyset, \{X\}$
- Arms: $do(\emptyset), do(X = 0), do(X = 1)$
- e.g., diet \rightarrow cholesterol \rightarrow health

SCM-MAB examples




- No non-manipulable variable
- Intervention sets: \emptyset , $\{X\}$, $\{Z\}$, $\{X, Z\}$
- Arms: $do(\emptyset)$, $do(X = 0)$, $do(X = 1)$, \dots , $do(X = 1, Z = 1)$

Structural Properties in SCM-MAB

Arms are **dependent** through underlying causal mechanism in SCM-MAB.

1. **Equivalence**: two arms share the **same** reward distribution




Structural Properties in SCM-MAB

Arms are **dependent** through underlying causal mechanism in SCM-MAB.

1. **Equivalence**: two arms share the **same** reward distribution

$$\mu_{x,z} = \mu_z$$




Structural Properties in SCM-MAB

Arms are **dependent** through underlying causal mechanism in SCM-MAB.

1. **Equivalence**: two arms share the **same** reward distribution

$$\mu_{x,z} = \mu_z$$

2. **Partial-orders**: one arm is always preferred to the other



Structural Properties in SCM-MAB


Arms are **dependent** through underlying causal mechanism in SCM-MAB.

1. **Equivalence**: two arms share the **same** reward distribution

$$\mu_{x,z} = \mu_z$$

2. **Partial-orders**: one arm is always preferred to the other

$$\mu_{x^*} \geq \mu_{z^*}$$



Structural Properties in SCM-MAB

Arms are **dependent** through underlying causal mechanism in SCM-MAB.


1. **Equivalence**: two arms share the **same** reward distribution

$$\mu_{x,z} = \mu_z$$

2. **Partial-orders**: one arm is always preferred to the other

$$\mu_{x^*} \geq \mu_{z^*}$$

3. **Expressions**: infer one arm's reward distr. from other arms' samples.



Structural Properties in SCM-MAB

Arms are **dependent** through underlying causal mechanism in SCM-MAB.

1. **Equivalence**: two arms share the **same** reward distribution


$$\mu_{x,z} = \mu_z$$

2. **Partial-orders**: one arm is always preferred to the other

$$\mu_{x^*} \geq \mu_{z^*}$$


3. **Expressions**: infer one arm's reward distr. from other arms' samples.

$$P_x(y) = \sum_z P(z|x) \sum_{x'} P(y|z, x') P(x')$$



Structural Property 1: Equivalence


Consider a graph:



$\mu_{x,z} = \mu_z$ based on Rule 3 of *do*-calculus (Pearl, 2000), $(Y \perp\!\!\!\perp X | Z)_{\mathcal{G}_{\overline{X,Z}}}$

Structural Property 1: Equivalence

Consider a graph:




$\mu_{x,z} = \mu_z$ based on Rule 3 of *do*-calculus (Pearl, 2000), $(Y \perp\!\!\!\perp X | Z)_{\mathcal{G}_{\overline{X,Z}}}$

Implication: play $do(z)$ instead of $do(x, z)$!

Structural Property 1: Equivalence

Consider a graph:




$\mu_{x,z} = \mu_z$ based on Rule 3 of *do*-calculus (Pearl, 2000), $(Y \perp\!\!\!\perp X | Z)_{\mathcal{G}_{\overline{X,Z}}}$

Implication: play $do(z)$ instead of $do(x, z)$!

Find out sets of variables with **unique** rewards.

Structural Property 1: Equivalence

Consider a graph:




$\mu_{x,z} = \mu_z$ based on Rule 3 of *do*-calculus (Pearl, 2000), $(Y \perp\!\!\!\perp X | Z)_{\mathcal{G}_{\overline{X,Z}}}$

Implication: play $do(z)$ instead of $do(x, z)$!

Find out **minimal** sets of variables with **unique** rewards.

Structural Property 1: Equivalence

Consider a graph:



$\mu_{x,z} = \mu_z$ based on Rule 3 of *do*-calculus (Pearl, 2000), $(Y \perp\!\!\!\perp X | Z)_{\mathcal{G}_{\overline{X,Z}}}$


Implication: play $do(z)$ instead of $do(x, z)$!

Definition (Minimal Intervention Set (MIS))

Given $\langle \mathcal{G}, Y, \mathbf{N} \rangle$, a set of variables $\mathbf{X} \subseteq \mathbf{V} \setminus \{Y\} \setminus \mathbf{N}$ is said to be a *minimal intervention set* if there is no $\mathbf{X}' \subset \mathbf{X}$ such that $\mu_{\mathbf{x}'} = \mu_{\mathbf{x}}$ for every SCM conforming to \mathcal{G} where $\mathbf{x}' \in \mathfrak{X}_{\mathbf{X}'}$ that is consistent with \mathbf{x} .

Structural Property 2: Partial-orderedness

Consider a graph:




$$\mu_x = \sum_z \mu_z P(z|x) \leq \sum_z \mu_{z^*} P(z|x) = \mu_{z^*}.$$

Note that, there is no partial-order between \emptyset and μ_z .

Structural Property 2: Partial-orderedness

Consider a graph:




$$\mu_x = \sum_z \mu_z P(z|x) \leq \sum_z \mu_{z^*} P(z|x) = \mu_{z^*}.$$

Note that, there is no partial-order between \emptyset and μ_z .

Implication: play $do(z)$ is preferred to playing $do(x)$.

Structural Property 2: Partial-orderedness

Consider a graph:



$$\mu_x = \sum_z \mu_z P(z|x) \leq \sum_z \mu_{z^*} P(z|x) = \mu_{z^*}.$$


Note that, there is no partial-order between \emptyset and μ_z .

Implication: play $do(z)$ is preferred to playing $do(x)$.

Find out sets of variables that is not **dominated** by other sets.

Structural Property 2: Partial-orderedness

Consider a graph:



$$\mu_x = \sum_z \mu_z P(z|x) \leq \sum_z \mu_{z^*} P(z|x) = \mu_{z^*}.$$


Note that, there is no partial-order between \emptyset and μ_z .

Implication: play $do(z)$ is preferred to playing $do(x)$.

Find out **minimal** sets of variables that is not **dominated** by other sets.

Structural Property 2: Partial-orderedness

Consider a graph:



$$\mu_x = \sum_z \mu_z P(z|x) \leq \sum_z \mu_{z^*} P(z|x) = \mu_{z^*}.$$

Note that, there is no partial-order between \emptyset and μ_z .

Implication: play $do(z)$ is preferred to playing $do(x)$.

Definition (Possibly-Optimal Minimal Intervention Set (POMIS))


Given $\langle \mathcal{G}, Y, \mathbf{N} \rangle$, let $\mathbf{X} \in \text{MISs}$. If there exists a SCM conforming to \mathcal{G} st

$$\mu_{\mathbf{x}^*} > \forall \mathbf{w} \in \text{MISs} \setminus \{\mathbf{X}\} \mu_{\mathbf{w}^*},$$

then \mathbf{X} is a *possibly-optimal minimal intervention set* wrt $\langle \mathcal{G}, Y, \mathbf{N} \rangle$.

Structural Property 2: Partial-orderedness

Consider a graph:



Since $do(z)$ becomes impossible, $do(x)$ is **not** dominated by other arms. Note that, there is no partial-order between \emptyset and μ_x .

Implication: play $do(\emptyset)$ and $do(x)$.

Definition (Possibly-Optimal Minimal Intervention Set (POMIS))

Given $\langle \mathcal{G}, Y, \mathbf{N} \rangle$, let $\mathbf{X} \in \text{MISs}$. If there exists a SCM conforming to \mathcal{G} st

$$\mu_{\mathbf{x}^*} > \forall \mathbf{w} \in \text{MISs} \setminus \{\mathbf{X}\} \mu_{\mathbf{w}^*},$$

then \mathbf{X} is a *possibly-optimal minimal intervention set* wrt $\langle \mathcal{G}, Y, \mathbf{N} \rangle$.


Structural Property 3: Relating (POMISs) Arms

- Q: how are samples from $\{do(\mathbf{z})\}_{\mathbf{z} \in \text{POMIS}}$ related to $do(\mathbf{x})$?

Structural Property 3: Relating (POMISs) Arms


- Q: how can we express $P_{\mathbf{x}}(\mathbf{v}')$ with $\{P_{\mathbf{z}}\}_{\mathbf{z} \in \text{POMIS}}$?
- ID: $P_{\mathbf{x}}(\mathbf{v}')$ from $P(\mathbf{v})$ (SP, 2006)
- zID: $P_{\mathbf{x}}(\mathbf{v}')$ from $P_{\mathbf{z}'}(\mathbf{v})$ for $\mathbf{Z}' \subseteq \mathbf{Z}$ (BP, 2012)
- **z²ID**: $P_{\mathbf{x}}(\mathbf{v}')$ from a set of experiments (this paper)

Structural Property 3: Relating POMISs — an example



POMISs are \emptyset , $\{B\}$, and $\{C\}$.


Structural Property 3: Relating POMISs — an example



POMISs are \emptyset , $\{B\}$, and $\{C\}$.

Can we express $P(y)$ with $P_b(\mathbf{v})$ only?


Structural Property 3: Relating POMISs — an example



POMISs are \emptyset , $\{B\}$, and $\{C\}$.

Can we express $P_c(y)$ with $P_b(\mathbf{v})$ and/or $P(\mathbf{v})$?

Structural Property 3: Relating POMISs — an example



POMISs are \emptyset , $\{B\}$, and $\{C\}$.

$$P(y) = \sum_{a,b,c} P_b(c|a) P_c(a, b, y)$$

$$P_b(y) = \sum_{a,c} P(c|a, b) \sum_{b'} P(y|a, b', c) P(a, b')$$

$$P_c(y) = \sum_{a,b} P(y|a, b, c) P(a, b)$$

$$P_c(y) = \sum_a P_b(y|a, c) P_b(a)$$

SCM-MAB algorithms

Incorporating Structural Properties into MAB algos.

What we know,

- **POMIS**: all arms vs. possibly-optimal arms
- **expressions**: utilize samples from other arms

Incorporating Structural Properties into MAB algos.

What we know,

- **POMIS**: all arms vs. possibly-optimal arms
- **expressions**: utilize samples from other arms

Two algorithms we considered:

- **Thompson sampling**: posterior sample for expected reward
→ approximate 'posterior distribution' w/ all available data.
- **kl-UCB**: upper bounds computed for expected reward
→ adjust 'upper bound' by taking account samples from other arms.

SCM-MAB algorithm: modified TS

taking advantage of **POMIS** and **z^2ID** .

```
function  $z^2$ -TS( $\mathcal{G}, Y, \mathbf{N}, T$ )  
   $\mathbb{Z} \leftarrow \mathbb{P}_{\mathcal{G}, Y}^{\mathbf{N}}$   
   $\mathbf{A} \leftarrow \{\mathbf{x} \in \mathcal{X}_{\mathbf{X}} \mid \mathbf{X} \in \mathbb{Z}\}$   
   $\hat{\theta}_{\mathbf{x}} \leftarrow \{P_{\mathbf{x}}(y)\} \cup \{z^2ID(\mathcal{G}, y, \mathbf{x}, \mathbb{Z}')\}_{\mathbb{Z}' \subseteq \mathbb{Z} \setminus \{\mathbf{x}\}}$  for  $\mathbf{x} \in \mathbf{A}$   
   $\mathbf{D} \leftarrow \{D_{\mathbf{x}} = \emptyset\}_{\mathbf{x} \in \mathbf{A}}$   
  for  $t$  in  $1, \dots, T$  do  
    for  $\mathbf{x} \in \mathbf{A}$  do  
       $\hat{\theta}_{\mathbf{x}}, \hat{s}_{\mathbf{x}}^2 \leftarrow \text{bMVWA}(\mathbf{D}, \hat{\theta}_{\mathbf{x}})$   
      Find  $\hat{\alpha}_{\mathbf{x}}, \hat{\beta}_{\mathbf{x}}$  such that Beta( $\hat{\alpha}_{\mathbf{x}}, \hat{\beta}_{\mathbf{x}}$ ) matching  $\hat{\theta}_{\mathbf{x}}, \hat{s}_{\mathbf{x}}^2$   
       $\theta_{\mathbf{x}} \sim \text{Beta}(\hat{\alpha}_{\mathbf{x}}, \hat{\beta}_{\mathbf{x}})$   
       $\mathbf{x}' \leftarrow \arg \max_{\mathbf{x} \in \mathbf{A}} \theta_{\mathbf{x}}$   
      Sample  $\mathbf{v}$  by  $do(\mathbf{x}')$  and append  $\mathbf{v}$  to  $D_{\mathbf{x}'}$ 
```

SCM-MAB algorithm: modified kl-UCB

taking advantage of **POMIS** and **z²ID**.

function z²-KL-UCB($\mathcal{G}, Y, N, T, f \leftarrow \ln(t) + 3 \ln(\ln(t))$)

Initialize $\mathbb{Z}, \mathbf{A}, \{\hat{\theta}_{\mathbf{x}}\}_{\mathbf{x} \in \mathbf{A}}, \mathbf{D}$

($\forall \mathbf{x} \in \mathbf{A}$) Sample \mathbf{v} by $do(\mathbf{x})$, and append \mathbf{v} to $D_{\mathbf{x}}$

for t in $|\mathbf{A}|, \dots, T$ **do**

$\hat{\theta}_{\mathbf{x}}, \hat{s}_{\mathbf{x}}^2 \leftarrow \text{bMVWA}(\mathbf{D}, \hat{\theta}_{\mathbf{x}})$ **for** $\mathbf{x} \in \mathbf{A}$

$\hat{N}_{\mathbf{x}} \leftarrow \hat{\theta}_{\mathbf{x}}(1 - \hat{\theta}_{\mathbf{x}})/\hat{s}_{\mathbf{x}}^2; \quad \hat{t} \leftarrow \sum_{\mathbf{x}} \hat{N}_{\mathbf{x}}$

$\mu = \left\{ \sup \left\{ \mu \in [0, 1] : \text{KL}(\hat{\theta}_{\mathbf{x}}, \mu) \leq \frac{f(\hat{t})}{\hat{N}_{\mathbf{x}}} \right\} \right\}_{\mathbf{x} \in \mathbf{A}}$




$\mathbf{x}' \leftarrow \arg \max_{\mathbf{x} \in \mathbf{A}} \mu_{\mathbf{x}}$

Sample \mathbf{v} by $do(\mathbf{x}')$, and append \mathbf{v} to $D_{\mathbf{x}'}$

Empirical Evaluation

Experimental settings


- 4 strategies: **Brute-force**, **MIS**, **POMIS**, **POMIS+**
- 2 base MAB algorithms: Thompson sampling (TS), kl-UCB
- 3 SCM-MAB problems, binary \mathbf{V}



- **1000** simulations

Experimental results (average cumulative regret)


base algorithm: kl-UCB



Performance: **POMIS+** > **POMIS** \geq **MIS** \geq **Brute-force**

Experimental results (average cumulative regret)

base algorithm: TS




Performance: **POMIS+** > **POMIS** \geq **MIS** \geq **Brute-force**

Conclusions

Bandit Landscape

Dimensions: functional assumptions, bandit type, reward type, etc.
We generalized MABs into a *causal* dimension.



SCM-MAB: stochastic iid reward, nonparametric, bandit feedback

Conclusions

We studied how to take advantage of a known causal graph in bandit setting:

- introduced: **SCM-MAB** w/ non-manipulability constraints
- characterized: Possibly-Optimal Minimal Intervention Set (**POMIS**)
- devised: **z^2 ID** to connect arms
- designed: SCM-MAB algorithms: **z^2 -TS**, **z^2 -kl-UCB**

Conclusions

We studied how to take advantage of a known causal graph in bandit setting:

- introduced: **SCM-MAB** w/ non-manipulability constraints
- characterized: Possibly-Optimal Minimal Intervention Set (**POMIS**)
- devised: **z^2 ID** to connect arms
- designed: SCM-MAB algorithms: **z^2 -TS**, **z^2 -kl-UCB**

Mahalo!

(= thank you)